

FOXY AUDIT · BUILD WITH GEMINI XPRIZE

How we build, and what it's for

Fifty days ago this was just an idea between me and my teammate. Now it's real, and most days it still feels like we're racing more than following any plan.

The honest split: we decide what to build and why. That part doesn't get outsourced to anything. Whether raw AI conversations should ever leave a developer's own machine — we decided no, never, full stop. Which feature actually matters this week vs. which one just sounds cool but isn't ready. Who we reach out to and how we say it without overselling ourselves. Those are ours.

AI is where we get our speed. Claude has gone through our own code more times than I can count at this point, and it's caught real stuff — a database isolation rule that wouldn't actually have held under the account we connect with, a sales page that still talked about an old version of the product we'd already moved past, a dashboard panel that was literally just making up numbers instead of showing real ones. That last one was embarrassing when we found it. Codex and GPT-5.6 have built actual features directly in our repo too. We read every change before accepting it — not because we don't trust the tools, just because if we can't explain our own code when someone asks, we shouldn't be shipping it.

Two people, fifty days, and tools doing a lot of the heavy lifting. That's genuinely the shape of it right now.

Who does what

One of us lives in the engineering — the hashing, the chain, the stuff that has to be exactly right or the whole idea falls apart. The other is out front — our socials, DMing people in the security world, writing the actual messages, asking strangers for five honest minutes of their time. We run Foxy Audit's own Twitter/X and LinkedIn ourselves too. Every post, every reply, every cold message to someone we actually respect — we write it. No template, no ghostwriter.

Who this is actually for

Not hospitals or banks directly. We're building for teams like us — two or three people who built something real and then hit the wall: an enterprise buyer's compliance team asks "prove your AI doesn't leak data," and there's no good answer besides "trust us." That stall can be the difference between a small team making it another few months or not. If what we're building turns that into "here, go verify it yourself," we've done something that matters past just our own company.

Where this goes

We don't have employees yet. We don't have paying customers yet either. Saying that straight feels better than dressing it up. The next real hire, when we get there — probably someone security-minded who can help us finish the parts we already know aren't done. Running the whole thing live end to end for real. Actually red-teaming our own detection instead of just believing it works because it worked in

our own tests. We've already talked to people whose opinion genuinely means something — a security consultant, a red-team founder, a professor who's been doing this since before we were born, a whole community of builders. Not customers yet. Just people willing to look at what we made and tell us the truth. That's usually how it actually starts.

Everything about this product has changed at least once in fifty days — features, the dashboard, even which model grades an interaction. The one thing that hasn't: your real data never leaves your own machine. That part didn't get to be fast or messy like everything else. It's the only decision we never let ourselves rush.

Foxy Audit · foxyaudit.tech